

# Integrated Network Analysis

Jinxi\_Hu-48528608, Samarth\_Grover-38220463

2025-11-09

## 1. Environment and Dependencies

```
library(readr)
library(igraph)
library(RColorBrewer)
library(ggplot2)
library(reshape2)
library(scales)
library(gridExtra)
set.seed(48528608)
```

## 2. Raw Data Loading and Preprocessing (before removing isolated nodes)

```
# Load raw node and edge data
a_nodes <- read.csv("data/nodes.csv")
a_edges <- read.csv("data/edges.csv")

# Keep nodes with non-empty titles
a_nodes_clean <- subset(a_nodes, !(is.na(title) | trimws(title) == ""))

# Filter valid edges (source and target both in node set)
a_edges_clean <- subset(a_edges, source %in% a_nodes_clean$local_id &
                           target %in% a_nodes_clean$local_id)

cat("== RAW DATA ==\n")
```

```
## == RAW DATA ==
```

```
cat("Total nodes (including potential isolated):", nrow(a_nodes_clean), "\n")
```

```
## Total nodes (including potential isolated): 2610
```

```
cat("Valid citation edges:", nrow(a_edges_clean), "\n\n")
```

```
## Valid citation edges: 3757
```

```

# Build undirected base graph (for overview and isolated node detection)
graph_raw <- graph_from_data_frame(a_edges_clean, vertices = a_nodes_clean, directed = FALSE)
# Remove multi-edges and self-loops
graph_raw <- simplify(graph_raw, remove.multiple = TRUE, remove.loops = TRUE)

cat("Raw undirected graph nodes:", vcount(graph_raw), "\n")

## Raw undirected graph nodes: 2610

cat("Raw undirected graph edges:", ecount(graph_raw), "\n\n")

## Raw undirected graph edges: 3722

# Compute isolated nodes
deg_all_raw <- degree(graph_raw, mode = "all")
raw_isolated <- V(graph_raw)[deg_all_raw == 0]
cat("Isolated nodes count:", length(raw_isolated), " (", round(length(raw_isolated)/vcount(graph_raw)*100, 2), "%)")

## Isolated nodes count: 1043 ( 39.96 %)

```

**3. Connected Data After Removing Isolated Nodes (for directed advanced analysis)**

```

# Load cleaned connected node/edge data (already separated from isolated set)
nodes_connected <- read.csv("data/nodes_connected.csv")
edges_connected <- read.csv("data/edges_connected.csv")

cat("== CONNECTED DATA ==\n")

## == CONNECTED DATA ==

cat("Connected nodes:", nrow(nodes_connected), "\n")

## Connected nodes: 1582

cat("Connected edges:", nrow(edges_connected), "\n\n")

## Connected edges: 3757

# Build directed graph
graph_connected <- graph_from_data_frame(edges_connected, vertices = nodes_connected, directed = TRUE)
graph_connected <- simplify(graph_connected, remove.multiple = TRUE, remove.loops = TRUE)

cat("Directed graph nodes:", vcount(graph_connected), "\n")

## Directed graph nodes: 1582

```

```

cat("Directed graph edges:", ecount(graph_connected), "\n")

## Directed graph edges: 3730

cat("Network density:", round(edge_density(graph_connected), 6), "\n")

## Network density: 0.001491

cat("Is directed:", is_directed(graph_connected), " | Is weighted:", is_weighted(graph_connected), "\n\n")

## Is directed: TRUE | Is weighted: FALSE

```

#### 4. Global Degree Statistics and Distributions (directed connected graph)

```

all_deg <- degree(graph_connected, mode = "all")
in_deg <- degree(graph_connected, mode = "in")
out_deg <- degree(graph_connected, mode = "out")

cat("==== GLOBAL DEGREE STATISTICS ====\n")

## === GLOBAL DEGREE STATISTICS ===

cat("Mean total degree:", round(mean(all_deg), 2), " | Range:", min(all_deg), "--", max(all_deg), " | SD:", round(sd(all_deg), 2))
## Mean total degree: 4.72 | Range: 0 - 110 | SD: 6.74

cat("Mean in-degree:", round(mean(in_deg), 2), " | Range:", min(in_deg), "--", max(in_deg), " | SD:", round(sd(in_deg), 2))
## Mean in-degree: 2.36 | Range: 0 - 104 | SD: 6.2

cat("Mean out-degree:", round(mean(out_deg), 2), " | Range:", min(out_deg), "--", max(out_deg), " | SD:", round(sd(out_deg), 2))
## Mean out-degree: 2.36 | Range: 0 - 29 | SD: 2.72

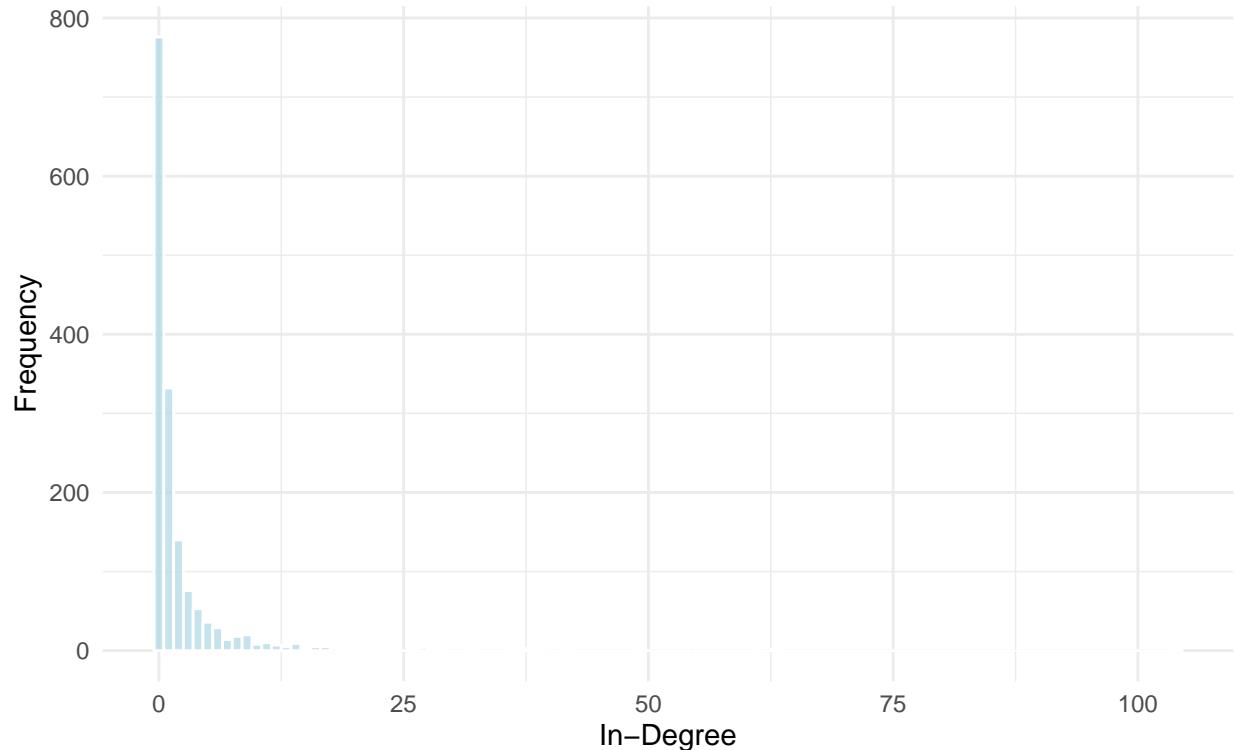
degree_data <- data.frame(
  node_id = V(graph_connected)$name,
  total_degree = all_deg,
  in_degree = in_deg,
  out_degree = out_deg,
  institution = V(graph_connected)$institution,
  subtopic = V(graph_connected)$subtopic,
  year = V(graph_connected)$year,
  citations = V(graph_connected)$citations
)

```

```
# In-degree distribution
ggplot(degree_data, aes(x = in_degree)) +
  geom_histogram(binwidth = 1, fill = "lightblue", alpha = 0.7, color = "white") +
  labs(title = "In-Degree Distribution", subtitle = paste("Connected nodes n=", vcount(graph_connected)),
       x = "In-Degree", y = "Frequency") + theme_minimal()
```

## In-Degree Distribution

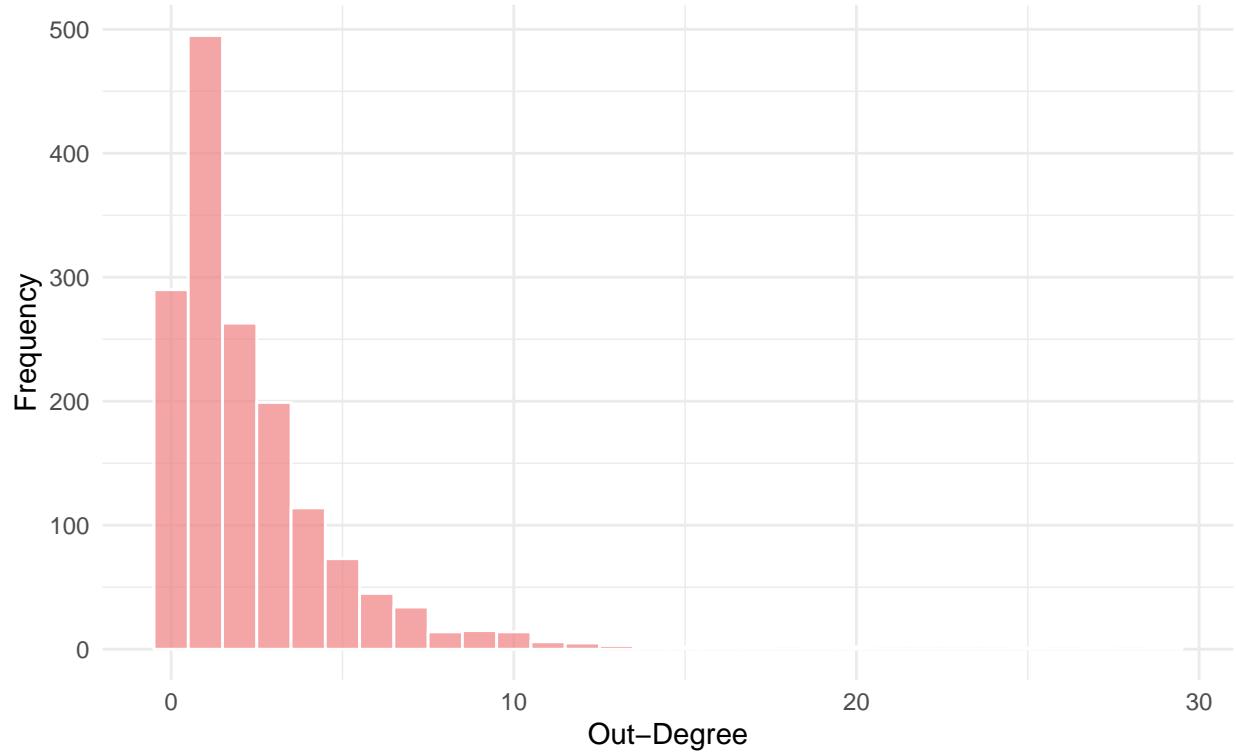
Connected nodes n= 1582



```
# Out-degree distribution
ggplot(degree_data, aes(x = out_degree)) +
  geom_histogram(binwidth = 1, fill = "lightcoral", alpha = 0.7, color = "white") +
  labs(title = "Out-Degree Distribution", subtitle = paste("Connected nodes n=", vcount(graph_connected)),
       x = "Out-Degree", y = "Frequency") + theme_minimal()
```

## Out-Degree Distribution

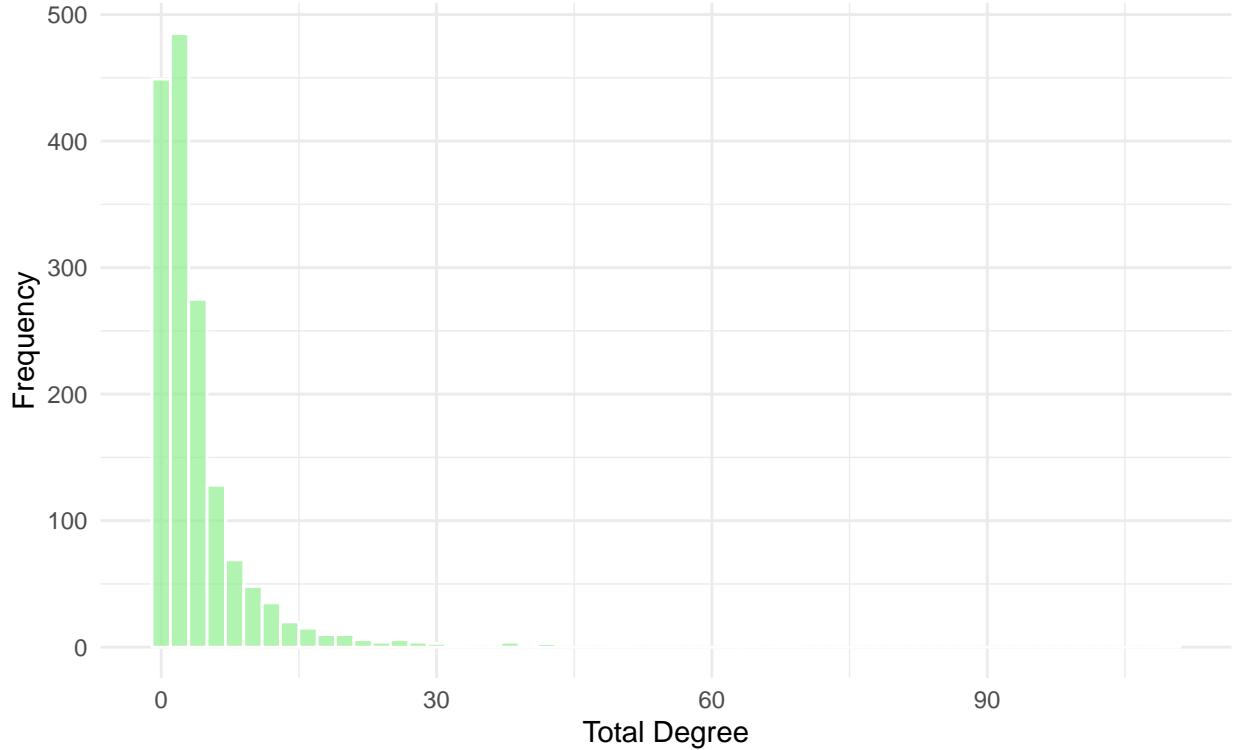
Connected nodes n= 1582



```
# Total-degree distribution
ggplot(degree_data, aes(x = total_degree)) +
  geom_histogram(binwidth = 2, fill = "lightgreen", alpha = 0.7, color = "white") +
  labs(title = "Total Degree Distribution", subtitle = paste("Connected nodes n=", vcount(graph_connect),
    x = "Total Degree", y = "Frequency") + theme_minimal()
```

## Total Degree Distribution

Connected nodes n= 1582



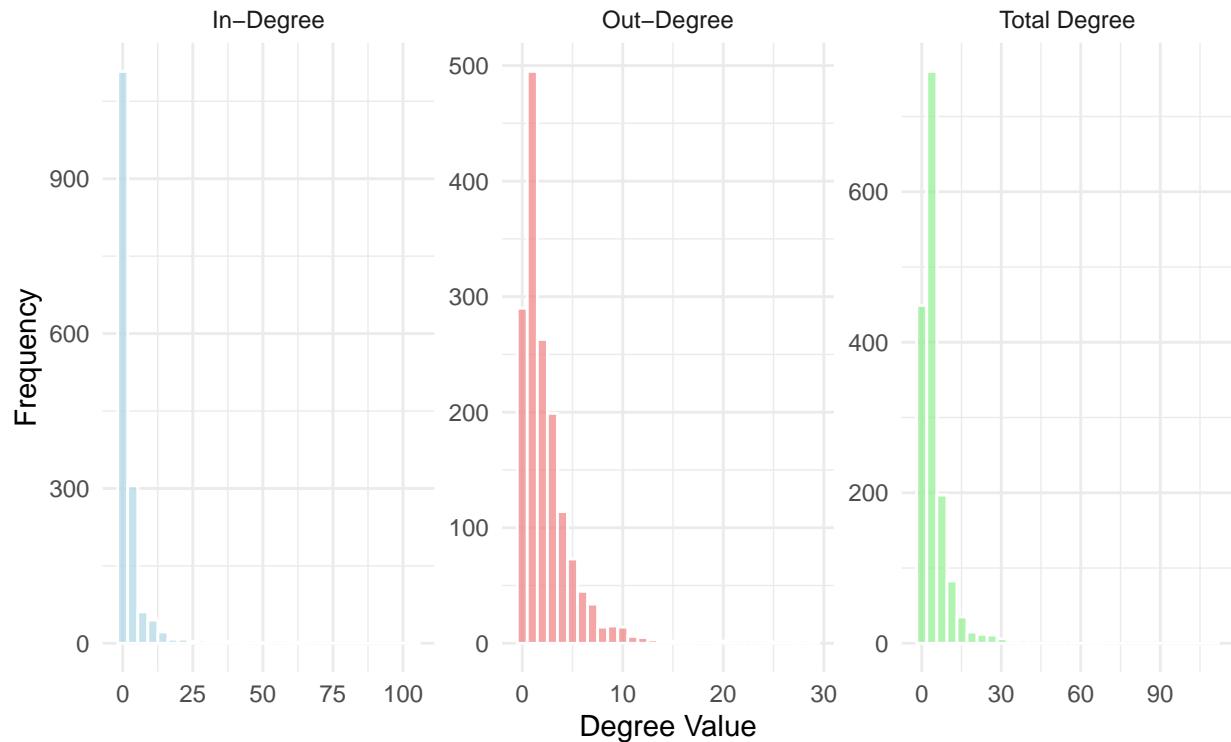
```
degree_long <- melt(degree_data[, c("in_degree", "out_degree", "total_degree")],  
                      variable.name = "degree_type", value.name = "degree_value")
```

```
## No id variables; using all as measure variables
```

```
degree_long$degree_type <- factor(degree_long$degree_type,  
                                    levels = c("in_degree", "out_degree", "total_degree"),  
                                    labels = c("In-Degree", "Out-Degree", "Total Degree"))  
  
ggplot(degree_long, aes(x = degree_value, fill = degree_type)) +  
  geom_histogram(alpha = 0.7, color = "white", bins = 30) +  
  facet_wrap(~degree_type, scales = "free") +  
  scale_fill_manual(values = c("In-Degree" = "lightblue", "Out-Degree" = "lightcoral", "Total Degree" = "#FFCCBC")) +  
  labs(title = "Degree Distribution Comparison", subtitle = "In-Degree vs Out-Degree vs Total Degree",  
       theme_minimal() + theme(legend.position = "none")
```

## Degree Distribution Comparison

### In-Degree vs Out-Degree vs Total Degree



## 5. Component Analysis (directed graph)

```
cat("==> COMPONENT ANALYSIS (DIRECTED GRAPH) ==>\n")

## ==> COMPONENT ANALYSIS (DIRECTED GRAPH) ==>

weak_comp <- components(graph_connected, mode = "weak")
strong_comp <- components(graph_connected, mode = "strong")

cat("Weakly connected components:", weak_comp$no, " | Largest weak component size:", max(weak_comp$csize))

## Weakly connected components: 58 | Largest weak component size: 1433 ( 90.58 %)

cat("Strongly connected components:", strong_comp$no, " | Largest strong component size:", max(strong_comp$csize))

## Strongly connected components: 1569 | Largest strong component size: 6 ( 0.38 %)

cat("Weak component sizes (Top 10):\n")

## Weak component sizes (Top 10):
```

```

print(head(sort(weak_comp$csize, decreasing = TRUE), 10))

## [1] 1433   27    9    5    5    4    3    3    3    3

cat("\nStrong component sizes (Top 10):\n")

## 
## Strong component sizes (Top 10):

print(head(sort(strong_comp$csize, decreasing = TRUE), 10))

## [1] 6 2 2 2 2 2 2 2 2 1

largest_nodes <- which(weak_comp$membership == which.max(weak_comp$csize))
largest_component <- induced_subgraph(graph_connected, largest_nodes)
cat("\n==== LARGEST WEAK COMPONENT STATISTICS ===\n")

## 
## === LARGEST WEAK COMPONENT STATISTICS ===

cat("Nodes:", vcount(largest_component), " | Edges:", ecount(largest_component), "\n")

## Nodes: 1433 | Edges: 3634

cat("Density:", round(edge_density(largest_component), 6), "\n")

## Density: 0.001771

cat("Average in-degree:", round(mean(degree(largest_component, mode = "in")), 2), " | Average out-degree:")

## Average in-degree: 2.54 | Average out-degree: 2.54

cat("Average total degree:", round(mean(degree(largest_component, mode = "all")), 2), "\n")

## Average total degree: 5.07

cat("Diameter:", diameter(largest_component, directed = FALSE), " | Average path length:", round(mean_diameter(largest_component), 2))

## Diameter: 13 | Average path length: 4.5

```

## 6. Centrality Analysis of Largest Component

```

cat("==== CENTRALITY (LARGEST COMPONENT) ====\n")

## === CENTRALITY (LARGEST COMPONENT) ====

```

```

cat("Calculating centralities... \n")

## Calculating centralities...

c_deg_in <- degree(largest_component, mode = "in", normalized = TRUE)
c_deg_out <- degree(largest_component, mode = "out", normalized = TRUE)
c_deg_all <- degree(largest_component, mode = "all", normalized = TRUE)
c_between <- betweenness(largest_component, directed = TRUE, normalized = TRUE)
c_close_in <- closeness(largest_component, mode = "in", normalized = TRUE)
c_close_out <- closeness(largest_component, mode = "out", normalized = TRUE)
c_eigen <- eigen_centrality(largest_component, directed = TRUE, scale = TRUE)$vector

## Warning: The `scale` argument of `eigen_centrality()` is deprecated as of igraph 2.1.1.
## i eigen_centrality() will always behave as if scale=TRUE were used.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.

c_pagerank <- page_rank(largest_component, directed = TRUE)$vector
c_hits_hub <- hub_score(largest_component)$vector

## Warning: `hub_score()` was deprecated in igraph 2.0.3.
## i Please use `hits_scores()` instead.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.

c_hits_auth <- authority_score(largest_component)$vector

## Warning: `authority_score()` was deprecated in igraph 2.1.0.
## i Please use `hits_scores()` instead.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.

centrality_summary <- data.frame(
  node_id = V(largest_component)$name,
  degree_in = c_deg_in,
  degree_out = c_deg_out,
  degree_all = c_deg_all,
  betweenness = c_between,
  closeness_in = c_close_in,
  closeness_out = c_close_out,
  eigenvector = c_eigen,
  pagerank = c_pagerank,
  hub = c_hits_hub,
  authority = c_hits_auth
)

cat("Top 5 In-Degree Centrality:\n"); print(head(centrality_summary[order(centrality_summary$degree_in,

```

```

## Top 5 In-Degree Centrality:

##      node_id degree_in
## P0060    P0060 0.07262570
## P0001    P0001 0.04399441
## P0009    P0009 0.04399441
## P0002    P0002 0.04259777
## P0199    P0199 0.04259777

cat("\nTop 5 PageRank:\n"); print(head(centrality_summary[order(centrality_summary$pagerank, decreasing=TRUE)]))

## 
## Top 5 PageRank:

##      node_id pagerank
## P0199    P0199 0.04046559
## P0031    P0031 0.03232105
## P0741    P0741 0.01974340
## P0116    P0116 0.01746391
## P0535    P0535 0.01572970

cat("\nTop 5 Betweenness Centrality:\n"); print(head(centrality_summary[order(centrality_summary$betweenness)]))

## 
## Top 5 Betweenness Centrality:

##      node_id betweenness
## P0011    P0011 0.002580361
## P0060    P0060 0.002094395
## P0003    P0003 0.001932865
## P0002    P0002 0.001732924
## P0001    P0001 0.001339179

cat("\nTop 5 Authority Score:\n"); print(head(centrality_summary[order(centrality_summary$authority, decreasing=TRUE)]))

## 
## Top 5 Authority Score:

##      node_id authority
## P0009    P0009 1.0000000
## P0060    P0060 0.9173024
## P0017    P0017 0.6726492
## P0014    P0014 0.6544079
## P0006    P0006 0.6283538

cat("==> Betweenness Centrality Distribution (Largest Component) ==>\n")

## ==> Betweenness Centrality Distribution (Largest Component) ==

```

```

bet_stats <- summary(c_between)
print(bet_stats)

##      Min. 1st Qu. Median     Mean 3rd Qu.    Max.
## 0.00e+00 0.00e+00 0.00e+00 2.12e-05 3.66e-06 2.58e-03

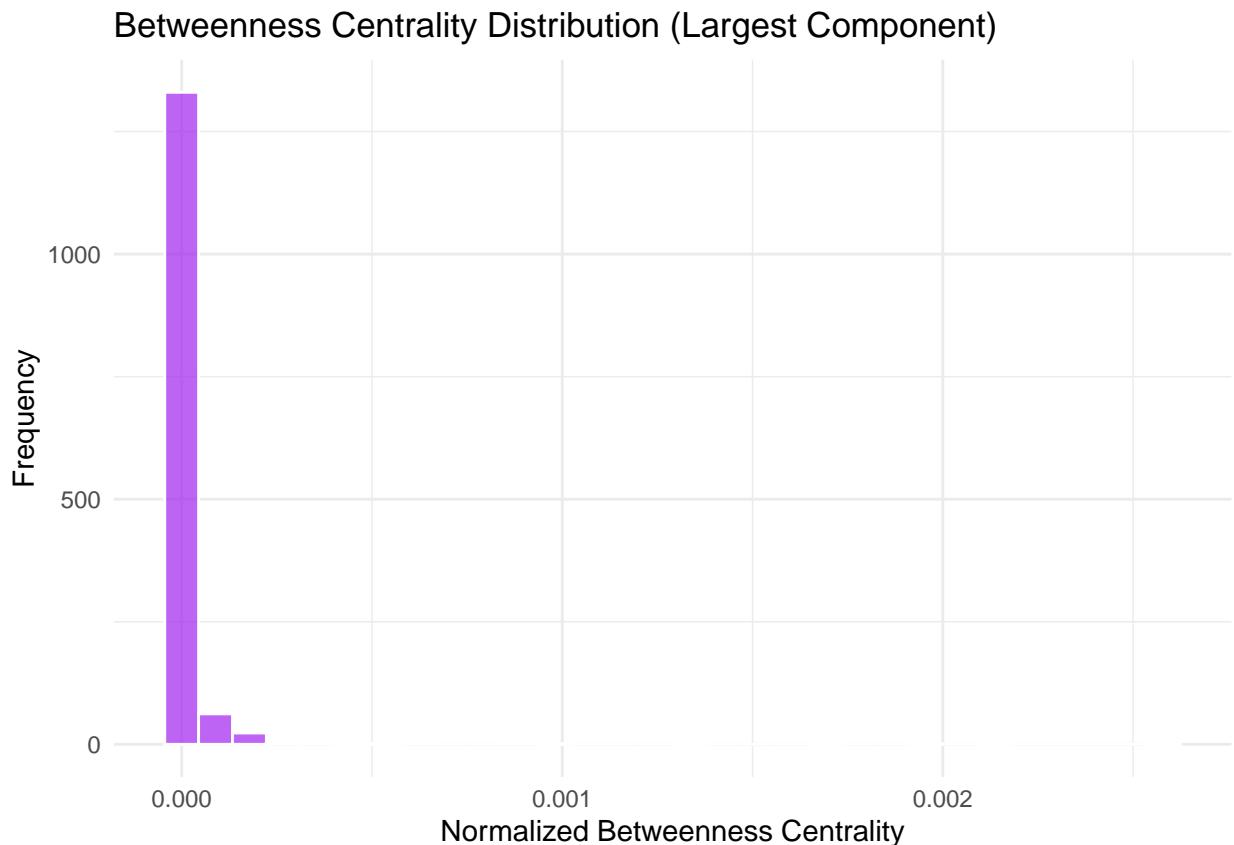
q95 <- quantile(c_between, 0.95)
high_between_ids <- which(c_between >= q95)
cat("Nodes in top 5% betweenness:", length(high_between_ids), "\n")

## Nodes in top 5% betweenness: 72

bet_df <- data.frame(node_id = V(largest_component)$name,
                      betweenness = c_between,
                      institution = V(largest_component)$institution,
                      subtopic = V(largest_component)$subtopic)

ggplot(bet_df, aes(x = betweenness)) +
  geom_histogram(bins = 30, fill = "purple", alpha = 0.7, color = "white") +
  labs(title = "Betweenness Centrality Distribution (Largest Component)", x = "Normalized Betweenness Centrality")

```



```
cat("\nTop 10 bridging nodes (betweenness):\n")
```

```

## 
## Top 10 bridging nodes (betweenness):

print(head(bet_df[order(bet_df$betweenness, decreasing = TRUE), c("node_id","betweenness","subtopic","institution")])

##          node_id      betweenness           subtopic
## P0011    P0011 0.0025803608 Machine Learning in Healthcare
## P0060    P0060 0.0020943953 Artificial Intelligence in Healthcare and Education
## P0003    P0003 0.0019328651 Artificial Intelligence in Healthcare and Education
## P0002    P0002 0.0017329242 Artificial Intelligence in Healthcare and Education
## P0001    P0001 0.0013391788 Machine Learning in Healthcare
## P0192    P0192 0.0009925863 Artificial Intelligence in Healthcare and Education
## P0484    P0484 0.0008632584 Artificial Intelligence in Healthcare and Education
## P0054    P0054 0.0008526921 Artificial Intelligence in Healthcare and Education
## P0136    P0136 0.0008494446 Artificial Intelligence in Healthcare and Education
## P0021    P0021 0.0006611386 Artificial Intelligence in Healthcare and Education
##          institution
## P0011    Stanford University
## P0060    University Of Cambridge
## P0003    Yale University
## P0002    Harvard University
## P0001    Stanford University
## P0192    Mayo Clinic In Arizona
## P0484    Stanford Health Care
## P0054    Stanford University
## P0136    Stanford University
## P0021    Emory University

cat("== Edge Betweenness (Largest Component) ==\n")

## == Edge Betweenness (Largest Component) ==

edge_bet <- edge_betweenness(largest_component, directed = TRUE)
cat("Number of edges:", length(edge_bet), "\n")

## Number of edges: 3634

print(summary(edge_bet))

##      Min.    1st Qu.     Median      Mean    3rd Qu.      Max.
## 1.000    2.000    5.667   25.071   17.462  2536.235

idx_top_e <- order(edge_bet, decreasing = TRUE)[1:10]
edge_pairs <- get.edges(largest_component, idx_top_e)
cat("\nTop 10 bridging edges: \n")

## 
## Top 10 bridging edges:

```


```

for(i in seq_along(idx_top_e)){
  from_n <- V(largest_component)$name[edge_pairs[i,1]]
  to_n <- V(largest_component)$name[edge_pairs[i,2]]
  cat(sprintf("%d. %s -> %s (%.4f)\n", i, from_n, to_n, edge_bet[idx_top_e[i]]))
}

## 1. P0001 -> P0003 (2536.2345)
## 2. P0003 -> P0054 (2051.3298)
## 3. P0484 -> P0011 (1825.9821)
## 4. P0192 -> P0136 (1811.0000)
## 5. P0011 -> P0192 (1789.1667)
## 6. P0003 -> P0021 (1346.8333)
## 7. P0002 -> P0107 (1293.0250)
## 8. P0060 -> P0001 (1211.1667)
## 9. P0060 -> P0002 (1039.8810)
## 10. P0011 -> P0003 (1002.4452)

ggplot(data.frame(edge_betweenness = edge_bet), aes(x = edge_betweenness)) +
  geom_histogram(bins = 50, fill = "orange", alpha = 0.7, color = "white") +
  labs(title = "Edge Betweenness Distribution (Largest Component)", x = "Edge Betweenness", y = "Frequency")

```



## 7. Community Detection Comparison (Largest Component)

```
cat("==> COMMUNITY DETECTION (Louvain vs Edge Betweenness) ==>\n")

## ==> COMMUNITY DETECTION (Louvain vs Edge Betweenness) ==>

largest_undirected <- as.undirected(largest_component, mode = "collapse")

## Warning: `as.undirected()` was deprecated in igraph 2.1.0.
## i Please use `as_undirected()` instead.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.

louvain_comm <- cluster_louvain(largest_undirected)
edgebet_comm <- cluster_edge_betweenness(largest_undirected, directed = FALSE)

cat("Louvain communities:", length(louvain_comm), " | Modularity:", round(modularity(louvain_comm), 4), "\n")

## Louvain communities: 21 | Modularity: 0.576

cat("Edge Betweenness communities:", length(edgebet_comm), " | Modularity:", round(modularity(edgebet_comm), 4), "\n")

## Edge Betweenness communities: 65 | Modularity: 0.5455

comm_compare <- data.frame(
  Method = c("Louvain", "Edge Betweenness"),
  Communities = c(length(louvain_comm), length(edgebet_comm)),
  Modularity = c(modularity(louvain_comm), modularity(edgebet_comm))
)

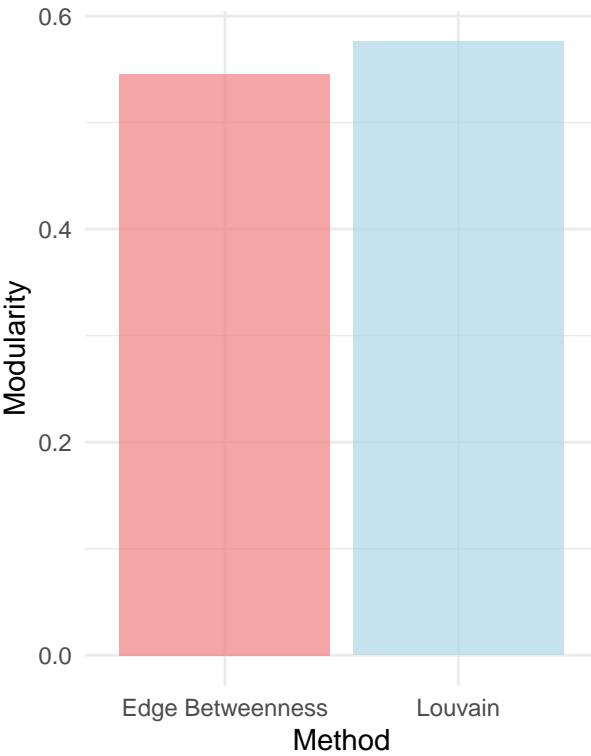
p1 <- ggplot(comm_compare, aes(x = Method, y = Communities, fill = Method)) +
  geom_bar(stat = "identity", alpha = 0.7) +
  scale_fill_manual(values = c("Louvain" = "lightblue", "Edge Betweenness" = "lightcoral")) +
  labs(title = "Number of Communities", y = "Count") + theme_minimal() + theme(legend.position = "none")

p2 <- ggplot(comm_compare, aes(x = Method, y = Modularity, fill = Method)) +
  geom_bar(stat = "identity", alpha = 0.7) +
  scale_fill_manual(values = c("Louvain" = "lightblue", "Edge Betweenness" = "lightcoral")) +
  labs(title = "Modularity Comparison", y = "Modularity") + theme_minimal() + theme(legend.position = "none")

grid.arrange(p1, p2, ncol = 2, top = "Community Detection Methods Comparison")
```

## Community Detection Methods Comparison

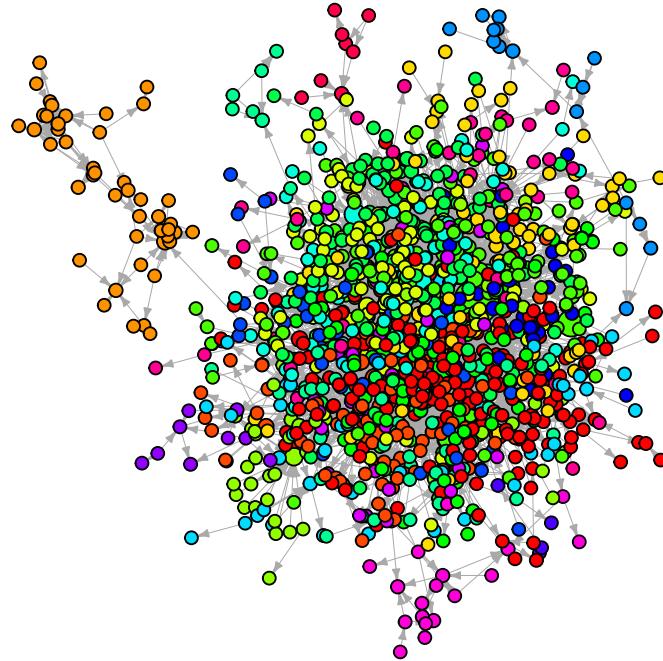
**Number of Communities**



```
# Louvain visualization + PageRank node sizing
page_rank_vals <- page_rank(largest_component, directed = TRUE)$vector
size_scaled <- 4 + 10 * (page_rank_vals - min(page_rank_vals)) / (max(page_rank_vals) - min(page_rank_vals))
comm_mem <- membership(louvain_comm)
if(length(louvain_comm) <= 12){
  comm_cols <- brewer.pal(max(3, length(louvain_comm)), "Set1")
} else {
  comm_cols <- rainbow(length(louvain_comm))
}
vertex_cols_comm <- comm_cols[comm_mem]
layout_large <- layout_with_fr(largest_component, niter = 800)

plot(largest_component, layout = layout_large,
      vertex.size = size_scaled,
      vertex.label = NA,
      edge.arrow.size = 0.3,
      edge.width = 0.5,
      vertex.color = vertex_cols_comm,
      main = "Largest Component: Louvain Communities + PageRank Node Size",
      sub = paste("Communities:", length(louvain_comm), " Modularity:", round(modularity(louvain_comm), 4)))
```

## Largest Component: Louvain Communities + PageRank Node Size



Communities: 21 Modularity: 0.576

## 8. Centrality Correlations

```
cat("== CENTRALITY CORRELATION MATRIX ==\n")
```

```
## == CENTRALITY CORRELATION MATRIX ==
```

```
cor_data <- centrality_summary[, c("degree_in", "degree_out", "betweenness", "closeness_in", "closeness_out")]
cor_mat <- cor(cor_data, use = "complete.obs")
print(round(cor_mat, 3))
```

	degree_in	degree_out	betweenness	closeness_in	closeness_out
## degree_in	1.000	0.002	0.616	-0.390	0.138
## degree_out	0.002	1.000	0.173	0.152	-0.262
## betweenness	0.616	0.173	1.000	-0.260	-0.101
## closeness_in	-0.390	0.152	-0.260	1.000	-0.277
## closeness_out	0.138	-0.262	-0.101	-0.277	1.000
## eigenvector	0.449	-0.068	0.458	-0.412	0.115
## pagerank	0.696	-0.057	0.460	-0.356	0.162
## authority	0.869	0.029	0.435	-0.311	0.125

```

## hub          0.096      0.670      0.110      0.036      -0.170
## eigenvector 0.449      0.696      0.869      0.096
## degree_in   0.449      0.696      0.869      0.096
## degree_out  -0.068     -0.057     0.029      0.670
## betweenness  0.458      0.460      0.435      0.110
## closeness_in -0.412     -0.356     -0.311     0.036
## closeness_out 0.115      0.162      0.125     -0.170
## eigenvector  1.000      0.701      0.331     -0.060
## pagerank    0.701      1.000      0.498     -0.012
## authority   0.331      0.498      1.000      0.187
## hub         -0.060     -0.012     0.187      1.000

```


```

cor_melt <- melt(cor_mat)

ggplot(cor_melt, aes(Var1, Var2, fill = value)) +
  geom_tile() +
  scale_fill_gradient2(low = "blue", high = "red", mid = "white", midpoint = 0, limit = c(-1,1), name =
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1)) +
  labs(title = "Centrality Correlations (Largest Component)", x = "", y = "") +
  geom_text(aes(label = round(value,2)), size = 3)

```

Centrality Correlations (Largest Component)



## 9. Isolated vs Connected Nodes Comparison

```
# Load isolated nodes data
nodes_isolated <- read.csv("data/nodes_isolated.csv")
cat("== ISOLATED NODE STATISTICS ==\n")

## == ISOLATED NODE STATISTICS ==

cat("Isolated nodes:", nrow(nodes_isolated), " | Year range:", min(nodes_isolated$year,na.rm=TRUE), "-"

## Isolated nodes: 1028 | Year range: 2015 - 2025

cat("Mean citations (isolated):", round(mean(nodes_isolated$citations, na.rm=TRUE),2), " | Mean references (isolated):"

## Mean citations (isolated): 9.97 | Mean references (isolated): 22.48

#
cat("== CONNECTED NODE STATISTICS ==\n")

## == CONNECTED NODE STATISTICS ==

cat("Connected nodes:", nrow(nodes_connected), " | Year range:", min(nodes_connected$year,na.rm=TRUE), "-"

## Connected nodes: 1582 | Year range: 2015 - 2025

cat("Mean citations (connected):", round(mean(nodes_connected$citations, na.rm=TRUE),2), " | Mean references (connected):"

## Mean citations (connected): 50.88 | Mean references (connected): 45.12

# Subtopic frequency comparison
sub_isolated <- table(nodes_isolated$subtopic)
sub_connected <- table(nodes_connected$subtopic)

sub_union <- unique(c(names(sub_isolated), names(sub_connected)))
comp_df <- data.frame(subtopic = sub_union,
                      isolated = as.numeric(sub_isolated[sub_union]),
                      connected = as.numeric(sub_connected[sub_union]))
comp_df$isolated[is.na(comp_df$isolated)] <- 0
comp_df$connected[is.na(comp_df$connected)] <- 0
comp_df$total <- comp_df$isolated + comp_df$connected
comp_df$isolated_prop <- ifelse(comp_df$total>0, comp_df$isolated/comp_df$total, NA)
comp_df_ord <- comp_df[order(comp_df$isolated_prop, decreasing = TRUE), ]

cat("Subtopics with highest isolation proportion (total>=5) Top 10:\n")

## Subtopics with highest isolation proportion (total>=5) Top 10:
```

```


print(head(comp_df_ord[comp_df_ord$total>=5, c("subtopic","isolated","connected","isolated_prop")],10))

##                                     subtopic isolated connected
## 7          Machine Learning and Algorithms      15        1
## 8          Machine Learning and Data Classification    34        4
## 3          Artificial Intelligence in Games       28        4
## 9          Machine Learning in Bioinformatics     26        7
## 12         Natural Language Processing Techniques   206       78
## 6          Artificial Intelligence in Law        44       23
## 4          Artificial Intelligence in Healthcare    30       18
## 11         Machine Learning in Materials Science    87       74
## 10         Machine Learning in Healthcare        111      167
## 5 Artificial Intelligence in Healthcare and Education  445      1206
##   isolated_prop
## 7      0.9375000
## 8      0.8947368
## 3      0.8750000
## 9      0.7878788
## 12     0.7253521
## 6      0.6567164
## 4      0.6250000
## 11     0.5403727
## 10     0.3992806
## 5      0.2695336

# Visualize top 15 subtopics by total count
comp_top15 <- head(comp_df[order(comp_df$total, decreasing = TRUE), ], 15)
plot_df <- melt(comp_top15[, c("subtopic","isolated","connected")], id.vars = "subtopic", variable.name

ggplot(plot_df, aes(x = reorder(subtopic, count), y = count, fill = type)) +
  geom_bar(stat = "identity") + coord_flip() +
  scale_fill_manual(values = c("isolated" = "lightcoral", "connected" = "lightblue"), labels = c("Isolated", "Connected"))
  labs(title = "Subtopics: Isolated vs Connected Nodes", subtitle = "Top 15 by total count", x = "Subtopic", y = "Count")
  theme_minimal() + theme(axis.text.y = element_text(size = 8))

```




```
# Institution frequency comparison
inst_iso <- table(nodes_isolated$institution)
inst_conn <- table(nodes_connected$institution)
inst_union <- unique(c(names(inst_iso), names(inst_conn)))
inst_df <- data.frame(institution = inst_union,
                      isolated = as.numeric(inst_iso[inst_union]),
                      connected = as.numeric(inst_conn[inst_union]))
inst_df$isolated[is.na(inst_df$isolated)] <- 0
inst_df$connected[is.na(inst_df$connected)] <- 0
inst_df$total <- inst_df$isolated + inst_df$connected

iso_top10 <- head(inst_df[order(inst_df$isolated, decreasing=TRUE), ], 10)
conn_top10 <- head(inst_df[order(inst_df$connected, decreasing=TRUE), ], 10)

iso_plot <- data.frame(
  institution = iso_top10$institution,
  frequency = iso_top10$isolated,
  type = "Isolated"
)
conn_plot <- data.frame(
  institution = conn_top10$institution,
  frequency = conn_top10$connected,
  type = "Connected"
)
inst_plot_df <- rbind(iso_plot, conn_plot)
```

```
ggplot(inst_plot_df, aes(x = reorder(institution, frequency), y = frequency, fill = type)) +
  geom_bar(stat = "identity") + coord_flip() + facet_wrap(~type, scales = "free") +
  scale_fill_manual(values = c("Isolated" = "lightblue", "Connected" = "lightsteelblue")) +
  labs(title = "Institutions: Isolated vs Connected Nodes (Top 10)", x = "Institution", y = "Count") +
  theme(axis.text.y = element_text(size = 7), legend.position = "none")
```



## 10. Overall Summary

```
cat("== SUMMARY ==\n")

## === SUMMARY ===

cat("Raw graph nodes:", vcount(graph_raw), " | Isolated nodes:", length(raw_isolated), " (", round(length(raw_isolated) / vcount(graph_raw) * 100, 2), "%)\n\n"
cat("Connected directed graph nodes:", vcount(graph_connected), " | Edges:", ecount(graph_connected), " | Density: ", density(graph_connected), "\n\n"
cat("Isolated nodes: ", length(raw_isolated), " | Density: ", density(raw_isolated), "\n\n")
cat("Connected components: ", length(components(graph_connected)), " | Max size: ", max_size(components(graph_connected)), "\n\n")
cat("Isolated components: ", length(components(raw_isolated)), " | Max size: ", max_size(components(raw_isolated)), "\n\n")
```

```
cat("Largest component nodes:", vcount(largest_component), " | Louvain communities:", length(louvain_com)

## Largest component nodes: 1433 | Louvain communities: 21 | Modularity: 0.576

cat("Highest PageRank node:", V(largest_component)$name[which.max(c_pageRank)], " (", round(max(c_pageR

## Highest PageRank node: P0199 ( 0.0405 )

cat("Highest betweenness node:", V(largest_component)$name[which.max(c_between)], " (", round(max(c_bett

## Highest betweenness node: P0011 ( 0.0026 )
```